

Three-Valued Bounded Model Checking with Cause-Guided Abstraction Refinement

Nils Timm*, Stefan Gruner, Dewald de Jager

Department of Computer Science, University of Pretoria, South Africa

Abstract

We present a technique for verifying concurrent software systems via SAT-based bounded model checking. It is based on a direct transfer of the system to be analysed and an LTL property into a formula that encodes the corresponding model checking problem. In our approach we first employ three-valued predicate abstraction. The state space of the resulting abstract system is then logically encoded, which saves us the expensive construction of an explicit state space model. The verification result can be obtained via two satisfiability checks. Our work includes the definition of the encoding and a theorem which states that the SAT result for an encoded verification task is equivalent to the result of the corresponding model checking problem. In case of an *unknown* result, the abstraction is automatically refined via our novel cause-guided refinement procedure that derives new predicates from the causes of uncertainty in the encoding. We also introduce an extension of the encoding by fairness constraints, which facilitates the verification of liveness properties. We have implemented our technique in an automatic verification tool that supports bounded LTL model checking under fairness.

Key words: Three-valued abstraction, Bounded model checking, Cause-guided abstraction refinement, Concurrent software systems, Fairness

1. Introduction

Three-valued abstraction (3VA) [1] is a well-established technique in software verification. It proceeds by generating an abstract state space model of the system to be analysed over the values *true*, *false* and *unknown*, where the latter value is used to represent the loss of information due to abstraction. For concurrent software systems composed of many processes, 3VA does not only replace concrete variables by predicates. It also abstracts away entire processes by summarising them into a single approximative component [2], which allows for a substantial reduction of the state space. The evaluation of temporal logic properties on models constructed via 3VA is known as *three-valued model checking* (3MC) [3]. In 3MC there exist three possible outcomes: *true* and *false* results

*Corresponding author

Email addresses: `ntimm@cs.up.ac.za` (Nils Timm), `sg@cs.up.ac.za` (Stefan Gruner)

Preprint submitted to Science of Computer Programming

December 12, 2017

1 can be immediately transferred to the modelled system, whereas an *unknown*
2 result reveals that abstraction refinement is necessary [4].

3 Verification techniques based on 3VA and 3MC typically assume that an
4 *explicit* three-valued state space model corresponding to the system to be anal-
5 ysed is constructed and explored [3]. However, explicit-state model checking is
6 known for its high memory demands in comparison to *symbolic* model checking
7 techniques like BDD-based model checking [5] and SAT-based bounded model
8 checking (BMC) [6]. The benefits of BMC are that its compressed state space
9 representation allows to handle larger systems than explicit-state techniques,
10 and that its performance profits from the advancements in the SAT solver tech-
11 nology. Although there exist a few works on *three-valued bounded model check-*
12 *ing*, these approaches are either solely defined for hardware systems [7], or they
13 require an explicit state space model as input which is then symbolically en-
14 coded [8]. It is however not efficient to first translate a given system into an
15 *explicit* state space model before encoding it symbolically for BMC.

16 In this paper we present and extend an approach to the verification of *con-*
17 *current software systems* based on an immediate transfer of the input system
18 and the property to be verified into a propositional logic formula that encodes
19 the corresponding bounded model checking problem [9]. Our approach first em-
20 ploys 3VA and thus profits from the state space reduction capabilities of this
21 technique. The state space of the resulting abstract system is then directly
22 encoded in propositional logic, which saves us the expensive construction of an
23 explicit state space model. Finally, the verification result, which can be *true*,
24 *false* or *unknown*, can be obtained via two satisfiability checks. An *unknown*
25 result indicates that the current abstraction is too coarse for a definite outcome.
26 While our previous work [9] did not provide an approach to the refinement of
27 three-valued abstractions, we additionally introduce a fully-automatic abstrac-
28 tion refinement procedure here. For this we enhance our encoding in the sense
29 that it comprises *causes of uncertainty*. Such causes can be missing information
30 about transitions or predicates. In case SAT-based three-valued model checking
31 yields *unknown*, our approach straightforwardly derives the associated causes of
32 uncertainty from unsatisfied clauses. The causes hint at additional predicates
33 over system variables that are required to rule out the uncertainty of the current
34 abstraction. We show that our novel iterative *cause-guided refinement* method
35 allows to automatically and quickly reach the right level of abstraction in order
36 to obtain a definite result in verification. Moreover, we report on promising
37 experimental results.

38 Our work includes the definition of the immediate encoding as well as a
39 proven theorem which states that the SAT result for an encoded verification
40 task is equivalent to the result of the corresponding model checking problem.
41 Furthermore, we introduce an extension of the encoding by *weak* and *strong*
42 *fairness* constraints, which facilitates the verification of liveness properties of
43 concurrent systems under realistic conditions. We have integrated the steps *ab-*
44 *straction*, *encoding*, *SAT solving* and *refinement* into a verification tool (avail-
45 able at www.github.com/ssfm-up/TVMC) that supports bounded LTL model
46 checking under fairness.

1 The remainder of this paper is organised as follows. In Section 2 we intro-
2 duce the systems that we consider in our software verification approach. Sec-
3 tion 3 provides the background on three-valued abstraction and bounded model
4 checking. Section 4 introduces our propositional logic encoding of software ver-
5 ification tasks and presents a theorem which states that the SAT result for an
6 encoded verification task is equivalent to the result of the corresponding model
7 checking problem. In Section 5 we show how our encoding can be augmented
8 with fairness constraints. Section 6 introduces our novel cause-guided abstrac-
9 tion refinement technique. In Section 7 we present the implementation of our
10 approach as well as experimental results. Section 8 discusses related work. We
11 conclude this paper in Section 9 and give an outlook on future work.

12 2. Concurrent Software Systems

13 We start with a brief introduction to the systems that we consider in our
14 work. A *concurrent software system* Sys consists of a number of possibly non-
15 uniform processes P_1 to P_n composed in parallel: $Sys = \parallel_{i=1}^n P_i$. It is defined
16 over a set of variables $Var = Var_s \cup \bigcup_{i=1}^n Var_i$ where Var_s is a set of shared
17 variables and Var_1, \dots, Var_n are sets of local variables associated with the pro-
18 cesses P_1, \dots, P_n , respectively. The state space over Var corresponds to the set
19 S_{Var} of all type-correct valuations of the variables. Given a state $s \in S_{Var}$ and
20 an expression e over Var , then $s(e)$ denotes the valuation of e in s . An example
21 for a concurrent system implementing mutual exclusion is depicted in Figure 1.

$$y : \text{semaphore where } y = 1;$$

$$P_1 :: \left[\begin{array}{l} \text{loop forever do} \\ \left[\begin{array}{l} 0: \text{acquire}(y, 1); \\ 1: \text{CRITICAL} \\ \text{release}(y, 1); \end{array} \right] \end{array} \right] \parallel P_2 :: \left[\begin{array}{l} \text{loop forever do} \\ \left[\begin{array}{l} 0: \text{acquire}(y, 1); \\ 1: \text{CRITICAL} \\ \text{release}(y, 1); \end{array} \right] \end{array} \right]$$

Figure 1: Concurrent system Sys .

22 Here we have two processes operating on a shared counting semaphore vari-
23 able y . Processes P_i can be formally represented as *control flow graphs* (CFGs)
24 $G_i = (Loc_i, \delta_i, \tau_i)$ where $Loc_i = \{[0]_2, \dots, [|Loc_i|]_2\}$ is a set of control loca-
25 tions given as binary numbers, $\delta_i \subseteq Loc_i \times Loc_i$ is a transition relation, and
26 $\tau_i : Loc_i \times Loc_i \rightarrow Op$ is a function labelling transitions with operations from a
27 set Op .

28 Definition 1 (Operations).

29 Let $Var = \{v_1, \dots, v_m\}$ be a set of variables. The set of operations Op on these
30 variables consists of all statements of the form $assume(e) : v_1 := e_1, \dots, v_m := e_m$
31 where e, e_1, \dots, e_m are expressions over Var .

32 Hence, every operation consists of a guard and a list of assignments. For
33 convenience, we sometimes just write e instead of $assume(e)$. Moreover, we
34 omit the *assume* part completely if the guard is *true*. The control flow graphs

₁ G_1 and G_2 corresponding to the processes of our example system are depicted in
₂ Figure 2. G_1 and G_2 also illustrate the semantics of the operations $acquire(y, 1)$
₃ and $release(y, 1)$.




Figure 2: Control flow graphs G_1 and G_2 composed in parallel.

₄ A concurrent system given by n individual control flow graphs G_1, \dots, G_n
₅ can be modelled by one composite CFG $G = (Loc, \delta, \tau)$ where $Loc = \times_{i=1}^n Loc_i$.
₆ G is the product graph of all individual CFGs. We assume that initially all pro-
₇ cesses of a concurrent system are at location 0. Moreover, we assume that a
₈ deterministic initialisation of the system variables is given by an assertion ϕ
₉ over Var . In our example we have that $\phi = (y = 1)$. Now, a computation
₁₀ of a concurrent system corresponds to a sequence where in each step one pro-
₁₁ cess is non-deterministically selected and the operation at its current location
₁₂ is attempted to be executed. In case the execution is not blocked by a guard,
₁₃ the variables are updated according to the assignment part and the process
₁₄ advances to the consequent control location. For verifying properties of concu-
₁₅ rent systems typically only *fair* computations where all processes infinitely often
₁₆ proceed are considered. We will discuss our notion of fairness in more detail in
₁₇ Section 5. The overall state space S of a concurrent system corresponds to the
₁₈ set of states over Var combined with the possible locations, i.e. $S = Loc \times S_{Var}$.
₁₉ Hence, each state in S is a tuple $\langle l, s \rangle$ with $l = (l_1, \dots, l_n) \in Loc$ and $s \in S_{Var}$.

₂₀ Control flow graphs allow to model concurrent systems formally. For an effi-
₂₁ cient verification it is additionally required to reduce the state space complexity.
₂₂ For this purpose, we use *three-valued predicate abstraction* [2]. Such an abstrac-
₂₃ tion is an approximation in the sense that all definite verification results (*true*,
₂₄ *false*) obtained for an abstract system can be transferred to the original sys-
₂₅ tem. Only *unknown* results necessitate abstraction refinement [4]. In abstract
₂₆ systems operations do not refer to concrete variables but to predicates $Pred =$
₂₇ $\{p_1, \dots, p_m\}$ over Var with the three-valued domain $\{true, unknown, false\}$. *Un-*
₂₈ *known*, typically abbreviated by \perp , is a valid truth value as we operate with the
₂₉ three-valued *Kleene logic* \mathcal{K}_3 [10] whose semantics is given by the truth tables
₃₀ in Figure 3.

\wedge	$true$	\perp	$false$	\vee	$true$	\perp	$false$	\neg	
$true$	$true$	\perp	$false$	$true$	$true$	$true$	$true$	$true$	$false$
\perp	\perp	\perp	$false$	\perp	$true$	\perp	\perp	\perp	\perp
$false$	$false$	$false$	$false$	$false$	$true$	\perp	$false$	$false$	$true$

Figure 3: Truth tables for the three-valued Kleene logic \mathcal{K}_3 .

Operations in abstract systems are of the following form:

$$assume(choice(a, b)) : p_1 := choice(a_1, b_1), \dots, p_m := choice(a_m, b_m)$$

- where $a, b, a_1, b_1, \dots, a_m, b_m$ are logical expressions over $Pred$ and $choice(a, b)$ -expressions have the following semantics:

Definition 2 (Choice Expressions).

Let s be a state over a set of three-valued predicates $Pred$. Moreover, let a and b be logical expressions over $Pred$. Then

$$s(choice(a, b)) = \begin{cases} true & \text{iff } s(a) \text{ is true,} \\ false & \text{iff } s(b) \text{ is true,} \\ \perp & \text{else.} \end{cases}$$

The application of three-valued predicate abstraction ensures that for any state s and for any expression $choice(a, b)$ in an abstract control flow graph the following holds: $s(a) = true \Rightarrow s(b) = false$ and $s(b) = true \Rightarrow s(a) = false$. In particular, this implies that $s(a)$ and $s(b)$ are never both $true$. Moreover, the following equivalences hold: $choice(true, false) \equiv true$, $choice(false, true) \equiv false$, $choice(false, false) \equiv \perp$, $choice(a, \neg a) \equiv a$, $choice(\neg a, a) \equiv \neg a$, as well as $choice(a, b) \equiv (a \vee \neg b) \wedge (a \vee b \vee \perp)$ and $\neg choice(a, b) \equiv choice(b, a)$.

A three-valued expression $choice(a, b)$ over $Pred$ approximates a Boolean expression e over Var , written $choice(a, b) \preceq e$, if and only if a logically implies e and b logically implies $\neg e$. The three-valued approximation relation can be straightforwardly extended to operations as described in [2]. An abstract system Sys' approximates a concrete system Sys , written $Sys' \preceq Sys$, if the systems have isomorphic CFGs and the operations in the abstract system approximate the corresponding ones in the concrete system. An example for an abstract system that approximates the concrete system in Figure 2 is depicted in Figure 4. For illustration: the abstract operation $(y > 0) := choice((y > 0), false)$ sets the predicate $(y > 0)$ to $true$ if $(y > 0)$ was $true$ before, and it never sets the predicate to $false$. This is a sound three-valued approximation of the concrete operation $y := y + 1$ over the predicate $(y > 0)$.

The state space of an abstract system is defined as $S = Loc \times S_{Pred}$ where S_{Pred} is the set of all possible valuations of the three-valued predicates in $Pred$.




Figure 4: Abstract system represented by control flow graphs G'_1 and G'_2 corresponding to the concrete control flow graphs G_1 and G_1 . Transitions are labelled with abstract operations over $Pred = \{(y > 0)\}$.

The state space corresponding to the abstraction of our example is thus $S =$

$$\begin{aligned}
& \{ \langle (0, 0), (y > 0) = \text{true} \rangle, \quad \langle (0, 0), (y > 0) = \perp \rangle, \quad \langle (0, 0), (y > 0) = \text{false} \rangle \\
& \langle (1, 0), (y > 0) = \text{true} \rangle, \quad \langle (1, 0), (y > 0) = \perp \rangle, \quad \langle (1, 0), (y > 0) = \text{false} \rangle \\
& \langle (0, 1), (y > 0) = \text{true} \rangle, \quad \langle (0, 1), (y > 0) = \perp \rangle, \quad \langle (0, 1), (y > 0) = \text{false} \rangle \\
& \langle (1, 1), (y > 0) = \text{true} \rangle, \quad \langle (1, 1), (y > 0) = \perp \rangle, \quad \langle (1, 1), (y > 0) = \text{false} \rangle \}.
\end{aligned}$$

So far we have seen how concurrent systems can be formally represented and abstracted. Next we will take a look on how model checking of abstracted systems is defined.

3. Three-Valued Bounded Model Checking

CFGs allow us to model the *control flow* of a concurrent system. The verification of a system additionally requires to explore a corresponding *state space* model. Since we use three-valued abstraction, we need a model that incorporates the truth values *true*, *false* and *unknown*. *Three-valued Kripke structures* are models with a three-valued domain for transitions and labellings of states:

Definition 3 (Three-Valued Kripke Structure).

A three-valued Kripke structure over a set of atomic predicates AP is a tuple $M = (S, \langle l^0, s^0 \rangle, R, L)$ where

- S is a finite set of states,
- $\langle l^0, s^0 \rangle \in S$ is the initial state,
- $R : S \times S \rightarrow \{\text{true}, \perp, \text{false}\}$ is a transition function with $\forall \langle l, s \rangle \in S : \exists \langle l', s' \rangle \in S : R(\langle l, s \rangle, \langle l', s' \rangle) \in \{\text{true}, \perp\}$,
- $L : S \times AP \rightarrow \{\text{true}, \perp, \text{false}\}$ is a labelling function that associates a truth value with each atomic predicate in each state.

A simple example for a three-valued Kripke structure M over $AP = \{p\}$ is depicted in Figure 5.




Figure 5: Three-valued Kripke structure.

1 A path π of a Kripke structure M is a sequence of states $\langle l^0, s^0 \rangle \langle l^1, s^1 \rangle \langle l^2, s^2 \rangle \dots$
2 with $R(\langle l^k, s^k \rangle, \langle l^{k+1}, s^{k+1} \rangle) \in \{\text{true}, \perp\}$. $\pi(k)$ denotes the k -th state of π ,
3 whereas π^k denotes the k -th suffix $\pi(k)\pi(k+1)\pi(k+2)\dots$ of π . By Π_M we de-
4 note the set of all paths of M starting in the initial state. Paths are considered
5 for the evaluation of temporal logic properties of Kripke structures.

6 As defined in [2], a concurrent system $Sys = \parallel_{i=1}^n P_i$ abstracted over a set of
7 predicates $Pred$ can be represented as a three-valued Kripke structure M over
8 $AP = Pred \cup \{(loc_i = l_i) \mid i \in [1..n], l_i \in Loc_i\}$ where the predicate $(loc_i = l_i)$
9 denotes that the process P_i is currently at control location l_i . The number of
10 states of a Kripke structure corresponding to a given system is exponential in the
11 number of its locations and variables. State explosion is the major challenge
12 in software model checking. One approach to cope with the state explosion
13 problem is to use a symbolic and therefore more compact representation of the
14 Kripke structure. In SAT-based bounded model checking [6] all possible path
15 prefixes up to a bound $b \in \mathbb{N}$ are encoded in a propositional logic formula. The
16 formula is then conjuncted with an encoding of the temporal logic property
17 to be checked. In case the overall formula is satisfiable, the satisfying truth
18 assignment characterises a witness path of length b for the property in the state
19 space of the encoded system. Hence, bounded model checking can be performed
20 via satisfiability solving. We now briefly recapitulate the syntax and bounded
21 semantics of the linear temporal logic (LTL):

Definition 4 (Syntax of LTL).

Let AP be a set of atomic predicates and $p \in AP$. The syntax of LTL formulae
 ψ is given by

$$\psi ::= p \mid \neg p \mid \psi \vee \psi \mid \psi \wedge \psi \mid \mathbf{G}\psi \mid \mathbf{F}\psi \mid \mathbf{X}\psi.$$

22 The temporal operator \mathbf{G} is read as *globally*, \mathbf{F} is read as *finally* (or *eventually*),
23 and \mathbf{X} is read as *next*. For the sake of simplicity, we omit the temporal
24 operator \mathbf{U} (*until*). Due to the extended domain of truth values in three-valued
25 Kripke structures, the bounded evaluation of LTL formulae is based on the

1 Kleene logic \mathcal{K}_3 (compare Section 2). Based on \mathcal{K}_3 , LTL formulae can be evalu-
 2 ated on b -bounded path prefixes of three-valued Kripke structures. Such finite
 3 prefixes $\pi(0) \dots \pi(b)$ can still represent infinite paths if the prefix has a *loop*,
 4 i.e. the last state $\pi(b)$ has a successor state that is also part of the prefix.

5 **Definition 5 (b -Loop).**

6 *Let π be a path of a three-valued Kripke structure M and let $r, b \in \mathbb{N}$ with $r \leq b$.
 7 Then π has a (b, r) -loop if $R(\pi(b), \pi(r)) \in \{\text{true}, \perp\}$ and π is of the form $v \cdot w^\omega$
 8 where $v = \pi(0) \dots \pi(r-1)$ and $w = \pi(r) \dots \pi(b)$. π has a b -loop if there exists
 9 an $r \in \mathbb{N}$ with $r \leq b$ such that π has a (b, r) -loop.*

10 For the bounded evaluation of LTL formulae on paths of Kripke structures
 11 we have to distinguish between paths *with* and *without* a b -loop.

12 **Definition 6 (Three-Valued Bounded Evaluation of LTL).**

13 *Let $M = (S, \langle l^0, s^0 \rangle, R, L)$ over AP be a three-valued Kripke structure. More-
 14 over, let $b \in \mathbb{N}$ and let π be a path of M with a b -loop. Then the b -bounded
 15 evaluation of an LTL formula ψ on π , written $[\pi \models_b^k \psi]$ where $k \leq b$ denotes
 16 the current position along the path, is inductively defined as follows:*

$$\begin{aligned}
 [\pi \models_b^k p] &\equiv L(\pi(k), p) \\
 [\pi \models_b^k \neg p] &\equiv \neg L(\pi(k), p) \\
 [\pi \models_b^k \psi \vee \psi'] &\equiv [\pi \models_b^k \psi] \vee [\pi \models_b^k \psi'] \\
 [\pi \models_b^k \psi \wedge \psi'] &\equiv [\pi \models_b^k \psi] \wedge [\pi \models_b^k \psi'] \\
 [\pi \models_b^k \mathbf{G}\psi] &\equiv \bigwedge_{k' \geq k} (R(\pi(k'), \pi(k' + 1)) \wedge [\pi \models_b^{k'} \psi]) \\
 [\pi \models_b^k \mathbf{F}\psi] &\equiv \bigvee_{k' \geq k} ([\pi \models_b^{k'} \psi] \wedge \bigwedge_{k''=k}^{k'-1} R(\pi(k''), \pi(k'' + 1))) \\
 [\pi \models_b^k \mathbf{X}\psi] &\equiv R(\pi(k), \pi(k + 1)) \wedge [\pi \models_b^{k+1} \psi]
 \end{aligned}$$

If π is a path without a b -loop then the b -bounded evaluation of ψ is defined as:

$$\begin{aligned}
 [\pi \models_b^k \mathbf{G}\psi] &\equiv \text{false} \\
 [\pi \models_b^k \mathbf{F}\psi] &\equiv \bigvee_{k'=k}^b ([\pi \models_b^{k'} \psi] \wedge \bigwedge_{k''=k}^{k'-1} R(\pi(k''), \pi(k'' + 1))) \\
 [\pi \models_b^k \mathbf{X}\psi] &\equiv \text{if } k < b \text{ then } R(\pi(k), \pi(k + 1)) \wedge [\pi \models_b^{k+1} \psi] \text{ else false}
 \end{aligned}$$

17 *The other cases are identical to the case where π has a b -loop. The universal*
 18 *bounded evaluation of ψ on an entire Kripke structure M is $[M \models_{U,b} \psi] \equiv$*
 19 *$\bigwedge_{\pi \in \Pi_M} [\pi \models_b^0 \psi]$. The existential bounded evaluation of ψ on a Kripke structure*
 20 *is $[M \models_{E,b} \psi] \equiv \bigvee_{\pi \in \Pi_M} [\pi \models_b^0 \psi]$.*

21 Checking temporal logic properties for three-valued Kripke structures is what
 22 is known as three-valued model checking [3]. Universal model checking can al-
 23 ways be transformed into existential model checking based on the equation
 24 $[M \models_{U,b} \psi] = \neg [M \models_{E,b} \neg \psi]$. From now on we only consider the existen-
 25 tial case, since it is the basis of satisfiability-based bounded model checking.

1 Bounded model checking [6] is typically performed incrementally, i.e. b is it-
 2 eratively increased until the property can be either proven or a completeness
 3 threshold [11] is reached. In the three-valued scenario there exist three possible
 4 outcomes: *true*, *false* and \perp . For our example Kripke structure M we have
 5 that $[M \models_{E,0} \mathbf{F}p]$ evaluates to \perp and $[M \models_{E,1} \mathbf{F}p]$ evaluates to *true*, which is
 6 witnessed by the 1-bounded path prefix $\langle l^0, s^0 \rangle \langle l^2, s^2 \rangle$.

7 It was shown in [2] that for a three-valued Kripke structure M modelling a
 8 concurrent system Sys abstracted over $Pred$ and an LTL formula ψ the following
 9 holds: $[M \models_{E,b} \psi] = \textit{true}$ implies that there exists an execution path of length
 10 b in Sys that satisfies ψ , and $[M \models_{E,b} \psi] = \textit{false}$ implies that no execution
 11 path of length b in Sys satisfies ψ . Hence, all definite model checking results
 12 obtained under three-valued abstraction can be immediately transferred to the
 13 concrete system Sys modelled by M , whereas an *unknown* result tells us that
 14 the current level of abstraction is too coarse.

15 In the next section we define a propositional logic encoding of three-valued
 16 bounded model checking tasks for abstracted concurrent systems. Our encod-
 17 ing allows to immediately transfer verification tasks into a propositional logic
 18 formulae that can be then processed via a SAT solver. Thus, the expensive
 19 construction of an explicit Kripke structure is not required in our approach.
 20 The state space of the system under consideration as well as the property to
 21 be checked will be implicitly contained in the propositional logic encoding, and
 22 the model checking result will be equivalent to the result of the corresponding
 23 satisfiability tests.

24 4. Propositional Logic Encoding

In our previous work [12] we showed that the three-valued bounded model
 checking problem $[M \models_{E,b} \psi]$, where M is given as an *explicit* Kripke structure,
 can be reduced to two classical SAT problems. Here we show that for a given
 system Sys abstracted over $Pred$, a temporal logic property ψ , and a bound
 $b \in \mathbb{N}$, it is not even necessary to consider the corresponding model checking
 problem. We can immediately construct a propositional logic formula $\llbracket Sys, \psi \rrbracket_b$
 such that:

$$[M \models_{E,b} \psi] = \begin{cases} \textit{true} & \textit{if} & \text{SAT}(\llbracket Sys, \psi \rrbracket_b[\perp \mapsto \textit{false}]) = \textit{true} \\ \textit{false} & \textit{if} & \text{SAT}(\llbracket Sys, \psi \rrbracket_b[\perp \mapsto \textit{true}]) = \textit{false} \\ \perp & \textit{else} \end{cases}$$

25 Here $[\perp \mapsto \textit{false}]$ resp. $[\perp \mapsto \textit{true}]$ denotes the mapping of all occurrences
 26 of \perp to *false* resp. *true*. Hence, it is not required to construct and explore an
 27 explicit Kripke structure M modelling the state space of Sys . All we need to do
 28 is to construct $\llbracket Sys, \psi \rrbracket_b$ and check its satisfiability in order to obtain the result
 29 of the corresponding model checking problem.

30 The formula $\llbracket Sys, \psi \rrbracket_b$ is defined over a set of Boolean atoms and the con-
 31 stants *true*, *false* and \perp . We now give a step-by-step description on how

1 $\llbracket Sys, \psi \rrbracket_b$ can be constructed for a concurrent system $Sys = \parallel_{i=1}^n P_i$ abstracted
2 over a set of predicates $Pred$ and given by a number of control flow graphs
3 $G_i = (Loc_i, \delta_i, \tau_i)$ with $1 \leq i \leq n$, a temporal logic property $\psi \in \text{LTL}$, and a
4 bound $b \in \mathbb{N}$. The construction of $\llbracket Sys, \psi \rrbracket_b$ is divided into the translation of
5 the abstract system into a formula $\llbracket Sys \rrbracket_b$ and the translation of the property
6 ψ into a formula $\llbracket \psi \rrbracket_b$.

We start with the encoding of the system, which first requires to encode its states as propositional logic formulae. Since a state of a concurrent system is a tuple $\langle l, s \rangle$ where l is a composite control flow location and s is a valuation of all predicates in $Pred$, we encode l and s separately. First, we introduce a set of Boolean atoms for the encoding of locations. A composite location $(l_1, \dots, l_n) \in Loc$ is a list of single locations $l_i \in Loc_i$ where $Loc_i = \{0, \dots, |Loc_i| - 1\}$ and i is the identifier of the associated process P_i . Each l_i is a binary number from the domain $\{[0]_2, \dots, [|Loc_i|]_2\}$. We assume that all these numbers have d_i digits where d_i is the number required to binary represent the maximum value $|Loc_i|$. We introduce the following set of Boolean atoms:

$$LocAtoms := \{l_i[j] \mid i \in [1..n], j \in [1..d_i]\}$$

7 Hence, for each process P_i of the system we introduce d_i Boolean atoms,
8 each referring to a distinct digit along the binary representation of its locations.
9 The atoms now allow us to define the following encoding of locations:

Definition 7 (Encoding of Locations).

Let the location $l_i \in \{0, \dots, |Loc_i| - 1\}$ be given as a binary number. Moreover, let $l_i(j)$ be a function evaluating to true if the j -th digit of l_i is 1, and to false otherwise. Then l_i can be encoded in propositional logic as follows:

$$enc(l_i) := \bigwedge_{j=1}^{d_i} ((l_i[j] \wedge l_i(j)) \vee (\neg l_i[j] \wedge \neg l_i(j)))$$

10 Let $l = (l_1, \dots, l_n)$ be a composite location. Then $enc(l) := \bigwedge_{i=1}^n enc(l_i)$.

11 Note that since the function $l_i(j)$ evaluates to *true* or *false* an encoding
12 $enc(l_i)$ can be always simplified to a conjunction of literals over $LocAtoms$. For
13 instance, the initial location $(0, 0)$ of our example system from Section 2 will be
14 encoded to $\neg l_1[1] \wedge \neg l_2[1]$ and the location $(0, 1)$ will be encoded to $\neg l_1[1] \wedge l_2[1]$.

Next, we encode the predicate part of states. Let $s \in S_{Pred}$ where $Pred = \{p_1, \dots, p_m\}$. We introduce the following set of Boolean atoms:

$$PredAtoms := \{p[j] \mid p \in Pred, j \in \{u, t\}\}$$

15 Hence, for each three-valued predicate p we introduce two Boolean atoms.
16 The atom $p[u]$ will let us indicate whether p evaluates to *unknown*, and $p[t]$ will
17 let us indicate whether it evaluates to *true* or *false*:

Definition 8 (Encoding of States over Predicates).

Let $p \in Pred$ and let $val \in \{true, \perp, false\}$. Then $(p = val)$ can be logically

encoded follows:

$$enc(p = val) := \begin{cases} \neg p[u] \wedge p[t] & \text{if } val = true \\ \neg p[u] \wedge \neg p[t] & \text{if } val = false \\ p[u] & \text{if } val = \perp \end{cases}$$

1 Let s be a state over $Pred$. Then $enc(s) := \bigwedge_{p \in Pred} enc(p = s(p))$.

2 For an overall state $\langle l, s \rangle \in S$ we consequently get $enc(\langle l, s \rangle) := enc(l) \wedge$
 3 $enc(s)$. Since $enc(\langle l, s \rangle)$ yields a conjunction of literals, there exists exactly one
 4 satisfying truth assignment $\alpha : LocAtoms \cup PredAtoms \rightarrow \{true, false\}$ for a
 5 state encoding. We denote the assignment characterising an encoded state $\langle l, s \rangle$
 6 by $\alpha_{\langle l, s \rangle}$. For instance, the initial state $\langle (0, 0), (y > 0) = true \rangle$ of our abstracted
 7 example system will be encoded to $Init = \neg l_1[1] \wedge \neg l_2[1] \wedge \neg p[u] \wedge p[t]$ where
 8 $p = (y > 0)$, i.e. we abbreviate $(y > 0)$ by p . The assignment characterising
 9 $Init$ is $\alpha_{\langle (0, 0), (y > 0) = true \rangle} : l_1[1] \mapsto false, l_2[1] \mapsto false, p[u] \mapsto false, p[t] \mapsto true$.

10 The encoding function enc can be extended to *logical expressions* in negation
 11 normal form (NNF), which we require for our later transition encoding:

Definition 9 (Encoding of Logical Expressions).

Let $p \in Pred$ and e, e' logical expressions in NNF over $Pred \cup \{true, \perp, false\}$.
 Let $val \in \{true, \perp, false\}$. Then the encoding of a logical expression is inductively defined as follows:

$$\begin{aligned} enc(val) &:= val \\ enc(\neg val) &:= \neg val \\ enc(p) &:= (p[u] \wedge \perp) \vee (\neg p[u] \wedge p[t]) \\ enc(\neg p) &:= (p[u] \wedge \perp) \vee (\neg p[u] \wedge \neg p[t]) \\ enc(e \wedge e') &:= enc(e) \wedge enc(e') \\ enc(e \vee e') &:= enc(e) \vee enc(e') \\ enc(choice(e, e')) &:= enc((e \vee NNF(\neg e')) \wedge (e \vee e' \vee \perp)) \end{aligned}$$

Next, we take a look at how the transition relation of an abstracted system can be encoded. We will construct a propositional logic formula

$$\llbracket Sys \rrbracket_b = Init_0 \wedge Trans_{0,1} \wedge \dots \wedge Trans_{b-1,b}$$

12 that exactly characterises path prefixes of length $b \in \mathbb{N}$ in the state space of the
 13 system Sys abstracted over $Pred$. Since we consider states as parts of such pre-
 14 fixes, we have to extend the encoding of states by index values $k \in \{0, \dots, b\}$
 15 where k denotes the position along a path prefix. For this we introduce the
 16 notion of indexed encodings. Let F be a propositional logic formula over
 17 $Atoms = LocAtoms \cup PredAtoms$ and the constants $true, false$ and \perp . Then
 18 F_k stands for $F[a/a_k \mid a \in Atoms]$. Our overall encoding will be thus defined
 19 over the set $Atoms_{[0,b]} = \{a_k \mid a \in Atoms, 0 \leq k \leq b\}$. An assignment $\alpha_{\langle l, s \rangle}$ to
 20 the atoms in a subset $Atoms_{[k,k]} \subseteq Atoms_{[0,b]}$ thus characterises a state $\langle l, s \rangle$ at
 21 position k of a path prefix, whereas an assignment $\alpha_{\langle l^0, s^0 \rangle \dots \langle l^b, s^b \rangle}$ to the atoms

1 in $Atoms_{[0,b]}$ characterises an entire path prefix $\langle l^0, s^0 \rangle \dots \langle l^b, s^b \rangle$. Since all ex-
 2 ecution paths start in the initial state of the system, we extend its encoding by
 3 the index 0, i.e. we get $Init_0 = \neg l_1[1]_0 \wedge \neg l_2[1]_0 \wedge \neg p[u]_0 \wedge p[t]_0$. The encoding of
 4 all possible state space transitions from position k to $k+1$ is defined as follows:

Definition 10 (Encoding of Transitions).

Let $Sys = \parallel_{i=1}^n P_i$ over $Pred$ be an abstracted concurrent system given by the
 single control flow graphs $G_i = (Loc_i, \delta_i, \tau_i)$ with $1 \leq i \leq n$. Then all possible
 transitions for position k to $k+1$ can be encoded in propositional logic as follows:

$$Trans_{k,k+1} :=$$

$$\bigvee_{i=1}^n \bigvee_{(l_i, l'_i) \in \delta_i} (enc(l_i)_k \wedge enc(l'_i)_{k+1} \wedge \bigwedge_{i' \neq i} (idle(i')_{k,k+1}) \wedge enc(\tau_i(l_i, l'_i))_{k,k+1})$$

where

$$idle(i')_{k,k+1} := \bigwedge_{j=1}^{d_{i'}} (l_{i'}[j]_k \leftrightarrow l_{i'}[j]_{k+1})$$

and

$$enc(\tau_i(l_i, l'_i))_{k,k+1} := \begin{aligned} & enc(choice(a, b))_k \\ & \wedge \bigwedge_{j=1}^m ((enc(a_j)_k \wedge enc(p_j = true)_{k+1}) \\ & \quad \vee (enc(b_j)_k \wedge enc(p_j = false)_{k+1}) \\ & \quad \vee (enc(\neg a_j \wedge \neg b_j)_k [\perp \mapsto true] \wedge enc(p_j = \perp)_{k+1})) \end{aligned}$$

5 assuming that $\tau_i(l_i, l'_i) = assume(choice(a, b)) : p_1 := choice(a_1, b_1), \dots, p_m :=$
 6 $choice(a_m, b_m)$.

7 Thus, we iterate over the system's processes P_i and over the processes' con-
 8 trol flow transitions $\delta_i(l_i, l'_i)$. Now we construct the k -indexed encoding of a
 9 source location l_i and conjunct it with the $(k+1)$ -indexed encoding of a desti-
 10 nation location l'_i . This gets conjuncted with the sub formula $\bigwedge_{i' \neq i} idle(i')_{k,k+1}$
 11 which encodes that all processes different to the currently considered process P_i
 12 are idle, i.e. do not change their control flow location, while P_i proceeds. The
 13 last part of the transition encoding concerns the operation associated with the
 14 control flow transition $\delta_i(l_i, l'_i)$: The sub formula $enc(\tau_i(l_i, l'_i))_{k,k+1}$ evaluates to
 15 $true$ for assignments $\alpha_{\langle l, s \rangle \langle l', s' \rangle}$ to the atoms in $Atoms_{[k,k+1]}$ that characterise
 16 pairs of states s and s' over $Pred$ where the guard of the operation $\tau_i(l_i, l'_i)$ is
 17 $true$ in s and the execution of the operation in s definitely results in the state
 18 s' . The operation encoding evaluates to \perp for states s and s' where the guard
 19 of the operation is \perp in s or where it is *unknown* whether the execution of
 20 the operation in s results in the state s' . In all other cases $enc(\tau_i(l_i, l'_i))_{k,k+1}$
 21 evaluates to $false$. Our transition encoding requires that an operation $\tau_i(l_i, l'_i)$
 22 assigns to all predicates in $Pred$: Thus, if a predicate p is not modified by the
 23 operation we assume that $p := p$ is part of the assignment list.

The encoding of the control flow transition $\delta_1(0, 1)$ of our abstract exam-
 ple system with $\tau_1(0, 1) = (assume(p) : p := choice(false, \neg p))$ (where p

abbreviates $(y > 0)$ yields the following:

$$\begin{aligned}
enc(0)_k &= \neg l_1[1]_k \\
\wedge \\
enc(1)_{k+1} &= l_1[1]_{k+1} \\
\wedge \\
idle(2)_{k,k+1} &= (l_2[1]_k \leftrightarrow l_2[1]_{k+1}) \\
\wedge \\
enc(\tau_1(0,1))_{k,k+1} &= ((p[u]_k \wedge \perp) \vee (\neg p[u]_k \wedge p[t]_k)) \wedge \\
& ((false \wedge (\neg p[u]_{k+1} \wedge p[t]_{k+1})) \\
& \vee (((p[u]_k \wedge \perp) \vee (\neg p[u]_k \wedge \neg p[t]_k)) \wedge (\neg p[u]_{k+1} \wedge \neg p[t]_{k+1})) \\
& \vee (((p[u]_k \wedge true) \vee (\neg p[u]_k \wedge p[t]_k)) \wedge (p[u]_{k+1})))
\end{aligned}$$

1 The encoding of the operation only evaluates to *true* for assignments to
2 the atoms in $Atoms_{[k,k+1]}$ that characterise a predicate state s at position k
3 with $s(p) = true$ and a state s' at position $k+1$ with $s'(p) = \perp$. An over-
4 all satisfying assignment for this encoding is $\alpha_{\langle(0,0),(y>0)=true\rangle\langle(1,0),(y>0)=\perp\rangle} :$
5 $l_1[1]_k \mapsto false, l_2[1]_k \mapsto false, l_1[1]_{k+1} \mapsto true, l_2[1]_{k+1} \mapsto false, p[u]_k \mapsto false,$
6 $p[t]_k \mapsto true, p[u]_{k+1} \mapsto true$ characterising the definite transition between
7 the pair of states $\langle(0,0),(y > 0) = true\rangle$ and $\langle(1,0),(y > 0) = \perp\rangle$. The as-
8 signments $\alpha_{\langle(0,l_2),(y>0)=true\rangle\langle(1,l_2),(y>0)=false\rangle}, \alpha_{\langle(0,l_2),(y>0)=\perp\rangle\langle(1,l_2),(y>0)=false\rangle},$
9 $\alpha_{\langle(0,l_2),(y>0)=\perp\rangle\langle(1,l_2),(y>0)=\perp\rangle}$ with $l_2 \in \{0,1\}$ yield *unknown* for the encoding
10 and hereby correctly characterise \perp -transitions in the abstract state space. All
11 other assignments yield *false* indicating that corresponding pairs of states do
12 not characterise valid transitions.

The encoding definitions now allow us to construct the propositional logic formula

$$\llbracket Sys \rrbracket_b = Init_0 \wedge Trans_{0,1} \wedge \dots \wedge Trans_{b-1,b}$$

13 that characterises all possible path prefixes of length $b \in \mathbb{N}$ in the state space
14 of the encoded system. Each assignment $\alpha : Atoms_{[0,b]} \rightarrow \{true, false\}$ that
15 satisfies the formula characterises a definite path prefix, whereas an assignment
16 that makes the formula evaluate to *unknown* characterises a prefix with some
17 \perp -transitions.

18 The second part of the encoding concerns the LTL property to be checked.
19 The three-valued bounded LTL encoding has been defined in [12] before. Here
20 we adjust it to our encodings of predicates and locations. Again, we distinguish
21 the cases where the property is evaluated on a path prefix with and without a
22 loop. The LTL encoding for the evaluation on prefixes with a loop is defined as:

Definition 11 (LTL Encoding with Loop).

Let p and $(loc_i = l_i) \in AP$, ψ and ψ' LTL formulae, and $b, k, r \in \mathbb{N}$ with $k, r \leq b$ where k is the current position, b the bound and r the destination position of the b -loop. Then the LTL encoding with a loop, ${}_r\llbracket \psi \rrbracket_b^k$, is defined as

follows:

$$\begin{aligned}
r[\llbracket loc_i = l_i \rrbracket_b^k] &\equiv enc(l_i)_k \\
r[\llbracket \neg(loc_i = l_i) \rrbracket_b^k] &\equiv \neg enc(l_i)_k \\
r[\llbracket p \rrbracket_b^k] &\equiv enc(p)_k \\
r[\llbracket \neg p \rrbracket_b^k] &\equiv enc(\neg p)_k \\
r[\llbracket \psi \vee \psi' \rrbracket_b^k] &\equiv r[\llbracket \psi \rrbracket_b^k] \vee r[\llbracket \psi' \rrbracket_b^k] \\
r[\llbracket \psi \wedge \psi' \rrbracket_b^k] &\equiv r[\llbracket \psi \rrbracket_b^k] \wedge r[\llbracket \psi' \rrbracket_b^k] \\
r[\llbracket \mathbf{G}\psi \rrbracket_b^k] &\equiv \bigwedge_{k'=\min(k,r)}^b r[\llbracket \psi \rrbracket_b^{k'}] \\
r[\llbracket \mathbf{F}\psi \rrbracket_b^k] &\equiv \bigvee_{k'=\min(k,r)}^b r[\llbracket \psi \rrbracket_b^{k'}] \\
r[\llbracket \mathbf{X}\psi \rrbracket_b^k] &\equiv r[\llbracket \psi \rrbracket_b^{succ(k)}]
\end{aligned}$$

1 where $succ(k) = k + 1$ if $k < b$ and $succ(k) = r$ else.

2 For a path prefix without a loop the LTL encoding is defined as:

3 **Definition 12 (LTL Encoding without Loop).**

4 Let ψ be an LTL formula and $b, k \in \mathbb{N}$ with $k \leq b$ where k is the current position
5 and b the bound. Then the LTL encoding without a loop, $\llbracket \psi \rrbracket_b^k$, is defined as
6 follows:

$$\begin{aligned}
\llbracket \mathbf{G}\psi \rrbracket_b^k &\equiv false \\
\llbracket \mathbf{F}\psi \rrbracket_b^k &\equiv \bigvee_{k'=k}^b \llbracket \psi \rrbracket_b^{k'} \\
\llbracket \mathbf{X}\psi \rrbracket_b^k &\equiv \text{if } k < b \text{ then } r[\llbracket \psi \rrbracket_b^{k+1}] \text{ else false}
\end{aligned}$$

7 The LTL encoding without a loop of the other cases is identical to the LTL
8 encoding with a loop.

An example encoding is $\llbracket \mathbf{F}p \rrbracket_2^0 = enc(p)_0 \vee enc(p)_1 \vee enc(p)_2$ which expresses that a predicate p holds *eventually*, i.e. at some position 0, 1 or 2 along a 2-prefix. Remember that a prefix $\langle l^0, s^0 \rangle \dots \langle l^b, s^b \rangle$ has a b -loop if there exists a transition from $\langle l^b, s^b \rangle$ to a previous state $\langle l^r, s^r \rangle$ along the prefix with $0 \leq r \leq b$. Hence, we can define a loop constraint based on our transition encoding: A prefix characterised by an assignment $\alpha_{\langle l^0, s^0 \rangle \dots \langle l^b, s^b \rangle}$ has definitely resp. maybe a b -loop if the loop constraint

$$\bigvee_{r=0}^b Trans_{b,r}$$

evaluates to *true* resp. *unknown* under $\alpha_{\langle l^0, s^0 \rangle \dots \langle l^b, s^b \rangle}$ where $Trans_{b,r}$ is defined according to Definition 10 but with k substituted by b and $k + 1$ by r . This now allows us to define the overall encoding of whether a concurrent system Sys satisfies an LTL formula ψ :

$$\llbracket Sys, \psi \rrbracket_b := \llbracket Sys \rrbracket_b \wedge \llbracket \psi \rrbracket_b$$

with

$$\llbracket \psi \rrbracket_b := \llbracket \psi \rrbracket_b^0 \vee \bigvee_{r=0}^b (Trans_{b,r} \wedge_r \llbracket \psi \rrbracket_b^0).$$

1 We have proven the following theorem that establishes the relation between
 2 the satisfiability result for $\llbracket Sys, \psi \rrbracket_b$ and the result of the corresponding model
 3 checking problem:

Theorem 1

Let M be a three-valued Kripke structure representing the state space of an abstracted concurrent system Sys , let ψ be an LTL formula and $b \in \mathbb{N}$. Then:

$$[M \models_{E,b} \psi] \equiv \begin{cases} true & \text{if } SAT(\llbracket Sys, \psi \rrbracket_b[\perp \mapsto false]) = true \\ false & \text{if } SAT(\llbracket Sys, \psi \rrbracket_b[\perp \mapsto true]) = false \\ \perp & \text{else} \end{cases}$$

4 *Proof.* See <http://www.cs.up.ac.za/cs/ntimm/ProofTheorem1.pdf>

5 Hence, via two satisfiability tests, one where \perp is mapped to *true* and one
 6 where it is mapped to *false*, we can determine the result of the corresponding
 7 model checking problem. Our encoding can be straightforwardly built based on
 8 the concurrent system, which saves us the expensive construction of an explicit
 9 state space model. In the next section we show that our encoding can be
 10 also easily augmented by fairness constraints, which allows us to check liveness
 11 properties of concurrent systems under realistic conditions.

12 **5. Extension to Fairness**

Our approach allows to check LTL properties of concurrent software systems via SAT solving. While the verification of safety properties like mutual exclusion does not require any fairness assumptions about the behaviour of the processes of the system, fairness is essential for verifying *liveness* properties under realistic conditions. The most common notions of fairness in verification are *unconditional*, *weak* and *strong* fairness: An unconditional fairness constraint claims that in an infinite computation, certain operations have to be infinitely often executed. A weak fairness constraint claims that in an infinite computation, each operation that is *continuously* enabled has to be infinitely often executed. A strong fairness constraint claims that in an infinite computation, each operation that is *infinitely often* enabled has to be infinitely often executed. All these types of constraints can be straightforwardly expressed in LTL. We now define these constraints for characterising fair, i.e. realistic, behaviour of our concurrent systems $Sys = \parallel_{i=1}^n P_i$ over $Pred$. Our unconditional fairness constraint is defined as:

$$ufair \equiv \bigwedge_{i=1}^n \bigvee_{(l_i, l'_i) \in \delta_i} \mathbf{GF}(executed(l_i, l'_i))$$

Hence, for each process some operation has to be executed infinitely often, i.e. each process proceeds infinitely often. Note that we model termination via a location with a self-loop. Thus, terminated processes can still proceed. The expression $executed(l_i, l'_i)$ can be easily defined in LTL. For this we extend the set $Pred$ by a progress predicate for each process: $Pred := Pred \cup \{progress_i \mid i \in [1..n]\}$. Moreover, we extend each operation as follows: $\tau_i(l_i, l'_i)$ sets $progress_i$ to *true* and all $progress_{i'}$ with $i' \neq i$ to *false*. Now $executed(l_i, l'_i)$ is defined as follows:

$$executed(l_i, l'_i) \equiv (loc_i = l_i) \wedge \mathbf{X}((loc_i = l'_i) \wedge progress_i).$$

Thus, an operation associated with a control flow transition (l_i, l'_i) is executed if $(loc_i = l_i)$ holds in the current state and $(loc_i = l'_i) \wedge progress_i$ holds in the next state. Next, we define our weak fairness constraint:

$$wfair \equiv \bigwedge_{i=1}^n \bigwedge_{(l_i, l'_i) \in \delta_i} (\mathbf{FG}(enabled(l_i, l'_i)) \rightarrow \mathbf{GF}(executed(l_i, l'_i)))$$

Hence, for each process, each continuously enabled operation has to be infinitely often executed. Instead of incorporating *each* operation in this type of constraint it is also possible to restrict the operations to crucial ones, which results in a shorter constraint and thus also restrains the complexity of model checking under fairness. For our running example it is for instance appropriate to just incorporate operations in $wfair$ that correspond to the successful acquisition of the semaphore. Note that $wfair$ can be easily transferred into negation normal form via the common propositional logic transformation rules such that it is conform with the definition of LTL. The expression $enabled(l_i, l'_i)$ can be defined as an LTL formula over locations and $Pred$ as follows:

$$enabled(l_i, l'_i) \equiv (loc_i = l_i) \wedge choice(a, b)$$

assuming that $\tau_i(l_i, l'_i) = assume(choice(a, b)) : p_1 := choice(a_1, b_1), \dots, p_m := choice(a_m, b_m)$. Thus, an operation associated with a control flow transition (l_i, l'_i) is enabled if $(loc_i = l_i)$ holds and the guard of the operation holds as well. Finally, we define our strong fairness constraint:

$$sfair \equiv \bigwedge_{i=1}^n \bigwedge_{(l_i, l'_i) \in \delta_i} (\mathbf{GF}(enabled(l_i, l'_i)) \rightarrow \mathbf{GF}(executed(l_i, l'_i)))$$

Hence, for each process, each operation that is enabled infinitely often has to be executed infinitely often. In model checking under fairness we can either check properties under specific constraints or we can combine all to a general one

$$fair \equiv wfair \wedge sfair.$$

Existential bounded model checking under fairness is now defined as:

$$[M \models_{E,b}^{fair} \psi] \equiv [M \models_{E,b} (fair \wedge \psi)]$$

Thus, we check whether there exists a b -bounded path that is fair and satisfies the property ψ . Such a model checking problem can be straightforwardly encoded in propositional logic based on our definitions in the previous section. We get

$$\llbracket Sys, fair \wedge \psi \rrbracket_b := \llbracket Sys \rrbracket_b \wedge \llbracket fair \wedge \psi \rrbracket_b,$$

1 which can be fed into a SAT solver in order to obtain the result of model checking
 2 ψ under fairness. Next, we introduce a systematic and fully-automatic approach
 3 to the refinement of three-valued abstractions in case the corresponding three-
 4 valued model checking problem yields *unknown*.

5 6. Cause-Guided Abstraction Refinement

6 In this section we present our approach to the refinement of three-valued ab-
 7 stractions in case the corresponding model checking problem yields an *unknown*
 8 result. Our abstractions represent uncertainty in the form of the constant \perp .
 9 SAT-based three-valued model checking is performed via two satisfiability tests,
 10 one where all occurrences of \perp are mapped to *true* in the propositional logic
 11 encoding $\llbracket Sys, \psi \rrbracket_b$ and one where its occurrences are mapped to *false*. Here we
 12 introduce an enhanced encoding that comprises the *causes of uncertainty*: Each
 13 \perp in the encoding gets superscripted with a *cause*, which can be missing infor-
 14 mation about a transition or a predicate. During the satisfiability tests all \perp 's
 15 are still treated the same, meaning that either all of them are mapped to *true* or
 16 all to *false* (compare Theorem 1). Once we have obtained an overall *unknown*
 17 model checking result, i.e. $\text{SAT}(\llbracket Sys, \psi \rrbracket_b[\perp \mapsto \text{true}]) = \text{true}$ for an assignment
 18 α and $\text{SAT}(\llbracket Sys, \psi \rrbracket_b[\perp \mapsto \text{false}]) = \text{false}$, we proceed as follows: We now have
 19 that the assignment α characterises an unconfirmed witness path for ψ contain-
 20 ing *unknowns*. Thus, this path is not present if all \perp 's get mapped to *false*. We
 21 determine the unsatisfied clauses of $\alpha(\llbracket Sys, \psi \rrbracket_b[\perp \mapsto \text{false}])$. All these clauses
 22 contain uncertainty in the sense of \perp 's and we will see that we can straightfor-
 23 wardly derive the corresponding causes. We then apply our novel *cause-guided*
 24 *abstraction refinement* which rules out the causes of uncertainty by adding new
 25 predicates to the abstraction. We will show that our fully-automatic iterative
 26 refinement approach enables us to quickly reach the right level of abstraction in
 27 order to obtain a definite model checking result.

28 The basis of our refinement technique is an enhanced encoding comprising
 29 causes of uncertainty:

Definition 13 (Causes of Uncertainty).

Let $\llbracket Sys, \psi \rrbracket_b$ be the propositional logic encoding of a three-valued bounded model checking problem corresponding to a concurrent system $Sys = \parallel_{i=1}^n P_i$ abstracted over $Pred$ where each process is given by a single control flow graph $G_i = (Loc_i, \delta_i, \tau_i)$ with $1 \leq i \leq n$. Uncertainty is represented in the encoding by the constant \perp . Each \perp in the encoding can be associated with a cause which

we define as follows:

$$cause \in \{p_k, (l_i, l'_i)_k\}$$

1 with $p \in Pred$, $0 \leq k \leq b$ and $l_i, l'_i \in Loc_i$.

2 We will use p_k in order to denote that the predicate p potentially evaluates
 3 to *unknown* at position k of the encoding, and with $(l_i, l'_i)_k$ we will denote that
 4 missing predicates over the guard of the operation $\tau_i(l_i, l'_i)$ potentially cause an
 5 *unknown* transition from position k to $k + 1$ in the encoding. Note that we
 6 refer to *potential* uncertainty in an encoding $\llbracket Sys, \psi \rrbracket_b$, since $\llbracket Sys, \psi \rrbracket_b$ always
 7 characterises *many* possible execution paths. For a *specific* path characterised
 8 by an assignment α to the atoms of $\llbracket Sys, \psi \rrbracket_b$ will see that we can refer to *actual*
 9 uncertainty. Next we show how causes of uncertainty can be integrated into the
 10 encoding in the sense of adding them as superscripts to the \perp 's. For this we
 11 introduce an enhanced encoding of abstract operations:

Definition 14 (Enhanced Encoding of Operations).

Let $Sys = \parallel_{i=1}^n P_i$ over $Pred$ be an abstracted concurrent system given by the single control flow graphs $G_i = (Loc_i, \delta_i, \tau_i)$ with $1 \leq i \leq n$. Then the encoding of abstract operations $\tau_i(l_i, l'_i) = \text{assume}(\text{choice}(a, b)) : p_1 := \text{choice}(a_1, b_1), \dots, p_m := \text{choice}(a_m, b_m)$ comprising the causes of uncertainty is defined as follows:

$$\begin{aligned} enc(\tau_i(l_i, l'_i))_{k,k+1} := & \quad enc(\text{choice}(a, b))_k \\ & \wedge \bigwedge_{j=1}^m ((enc(a_j)_k \wedge enc(p_j = \text{true})_{k+1}) \\ & \quad \vee (enc(b_j)_k \wedge enc(p_j = \text{false})_{k+1}) \\ & \quad \vee (enc(\neg a_j \wedge \neg b_j)_k [\perp / \text{true}] \wedge enc(p_j = \perp)_{k+1})) \end{aligned}$$

with

$$enc(\text{choice}(a, b))_k := enc((a \vee NNF(\neg b)) \wedge (a \vee b \vee \perp^{(\mathbf{l}, \mathbf{l}')_k}))_k$$

and the following inductive definition of the encoding of logical expressions e, e' over predicates $p \in Pred$.

$$\begin{aligned} enc(p)_k &:= (p[u]_k \wedge \perp^{\mathbf{p}_k}) \vee (\neg p[u]_k \wedge p[t]_k) \\ enc(\neg p)_k &:= (p[u]_k \wedge \perp^{\mathbf{p}_k}) \vee (\neg p[u]_k \wedge \neg p[t]_k) \\ enc(e \wedge e')_k &:= enc(e)_k \wedge enc(e')_k \\ enc(e \vee e')_k &:= enc(e)_k \vee enc(e')_k \end{aligned}$$

This definition enhances our previous encoding Definitions 9 and 10 in terms of superscripting each \perp with a corresponding cause. We will ignore the causes and treat all \perp 's the same during the satisfiability checks. Hence, the enhanced encoding is equivalent to the standard encoding. However, in case of an *unknown* model checking result, the causes will become crucial and will allow us to immediately derive expedient refinement steps. For illustration, we encode

the following abstract operation:

$$\tau(l_i, l'_i) = \text{assume}(\text{choice}(\text{false}, \text{false})) : p := \text{choice}(p, \neg p)$$

Remember that $\text{choice}(\text{false}, \text{false}) \equiv \perp$. For the enhanced encoding we get:

$$\begin{aligned} \text{enc}(\tau(l_i, l'_i))_{k,k+1} = \\ \perp^{(l_i, l'_i)_k} \wedge (((p[u]_k \wedge \perp^{\mathbf{pk}}) \vee (\neg p[u]_k \wedge p[t]_k)) \wedge \neg p[u]_{k+1} \wedge p[t]_{k+1}) \\ \vee (((p[u]_k \wedge \perp^{\mathbf{pk}}) \vee (\neg p[u]_k \wedge \neg p[t]_k)) \wedge \neg p[u]_{k+1} \wedge \neg p[t]_{k+1}) \end{aligned}$$

Thus, uncertainty may be caused by the unknown guard of the abstract operation $\tau(l_i, l'_i)$ or by the predicate p evaluating to \perp at position k . Actual uncertainty along a path characterised by an assignment α is only present if the \perp 's occur in clauses unsatisfied under $\alpha[\perp \mapsto \text{false}]$. Then we can utilise the causes attached to the \perp 's in order to rule out the uncertainty. We will now introduce our iterative abstraction-based model checking procedure with cause-guided refinement:

Procedure 1 (Iterative Abstraction-Based Model Checking).

Let $G = (Loc, \delta, \tau)$ be the concrete control flow graph representing a concurrent system Sys defined over a set of variables Var . Moreover, let ψ be an LTL formula to be checked for Sys . The corresponding bounded model checking problem can be solved via three-valued abstraction refinement and satisfiability solving as follows:

1. Initialise the set of predicates $Pred$ with the atomic propositions over Var referenced in ψ . Initialise the bound b with 1.
2. Construct the abstract control flow graph $G_a = (Loc_a, \delta_a, \tau_a)$ representing Sys abstracted over the current set $Pred$ via the abstractor 3Spot [2].
3. Encode the three-valued bounded model checking problem $[M(G_a) \models_{E,b} \psi]$ for the current bound b in propositional logic, which yields the formula $\llbracket Sys(G_a), \psi \rrbracket_b$.
4. Apply SAT-based three-valued bounded model checking (according to Theorem 1):
 - (a) If the result is $[M(G_a) \models_{E,b} \psi] = \text{true}$, then there exists a b -bounded witness path for ψ in the state space of Sys . The path is characterised by an assignment α satisfying $\llbracket Sys(G_a), \psi \rrbracket_b[\perp \mapsto \text{false}]$. Return α .
 - (b) If the result is $[M(G_a) \models_{E,b} \psi] = \text{false}$, then there does not exist a b -bounded witness path for ψ in the state space of Sys . Terminate if b has reached the completeness threshold [11] of the verification task. Otherwise increment b and go to 3.
 - (c) If the result is $[M(G_a) \models_{E,b} \psi] = \perp$, then it is *unknown* whether there exists a b -bounded witness path for ψ in the state space of Sys . An unconfirmed witness path for ψ with *unknowns* is characterised by an assignment α satisfying $\llbracket Sys(G_a), \psi \rrbracket_b[\perp \mapsto \text{true}]$ but not satisfying $\llbracket Sys(G_a), \psi \rrbracket_b[\perp \mapsto \text{false}]$. Apply the procedure *Cause-Guided Abstraction Refinement*, which updates $Pred$, and go to 2.

Procedure 2 (Cause-Guided Abstraction Refinement).

Let $G = (Loc, \delta, \tau)$ be the concrete composite control flow graph representing a concurrent system $Sys = \parallel_{i=1}^n P_i$ defined over a set of variables Var and let $G_i = (Loc_i, \delta_i, \tau_i)$ with $1 \leq i \leq n$ the corresponding single CFGs. Let ψ be an LTL formula to be checked for Sys . Moreover, let $G_a = (Loc_a, \delta_a, \tau_a)$ be the abstract composite control flow graph representing Sys abstracted over a set of predicates $Pred$ and let $[M(G_a) \models_{E,b} \psi]$ be the corresponding three-valued bounded model checking problem with $[M(G_a) \models_{E,b} \psi] = \perp$. Then for the propositional logic encoding $\llbracket Sys(G_a), \psi \rrbracket_b$ the following holds: $SAT(\llbracket Sys(G_a), \psi \rrbracket_b[\perp \mapsto true]) = true$ and the solver additionally returns a corresponding satisfying assignment α characterising an unconfirmed witness path. $SAT(\llbracket Sys(G_a), \psi \rrbracket_b[\perp \mapsto false]) = false$. Now the abstraction can be automatically refined by updating $Pred$ as follows:

1. Determine the set U of clauses of $\llbracket Sys, \psi \rrbracket_b$ that are unsatisfied under the assignment $\alpha[\perp \mapsto false]$. (Each $u \in U$ must contain at least one \perp since we have that $SAT(\llbracket Sys(G_a), \psi \rrbracket_b[\perp \mapsto true]) = true$.)
2. Determine a set $Causes$ such that for each $u \in U$ there exists a *cause* $\in Causes$ with \perp^{cause} is contained in u .
3. For each *cause* $\in Causes$:
 - (a) If *cause* $= (l_i, l'_i)_k$ with $l_i, l'_i \in Loc_i$, $1 \leq i \leq n$ and $0 \leq k \leq b$, then the value of the k -th transition along the unconfirmed witness path characterised by α is *unknown*. The transition from k to $k+1$ is associated with an operation $\tau_i(l_i, l'_i) = assume(e) : v_1 := e_1, \dots, v_m := e_m$ of the concrete control flow graph G_i . $\tau_i(l_i, l'_i)$ can be straightforwardly derived from G_i . Add the atomic propositions occurring in the assume condition e as new predicates to $Pred$.
 - (b) If *cause* $= p_k$ with $p \in Pred$ and $0 \leq k \leq b$, then the value of p is *unknown* at position k of the unconfirmed witness path characterised by α , i.e. $\alpha(p[u]_k) = true$. Let $k' < k$ be the last predecessor of k with $\alpha(p[u]_{k'}) = false$, i.e. the last position where the value of p is known. The transition from position k' to $k'+1$ is associated with an operation $\tau_i(l_i, l'_i) = assume(e) : v_1 := e_1, \dots, v_m := e_m$ of a concrete control flow graph G_i . Missing information about this concrete operation in terms of predicates is the cause of uncertainty in the current abstraction. $\tau_i(l_i, l'_i)$ can be straightforwardly derived based on G_i and $\alpha(LocAtoms_{[k', k'+1]})$, which indicates the corresponding control flow locations. Let $wp_{\tau_i(l_i, l'_i)}(p) = p[v_1/e_1, \dots, v_m/e_m]$ be the weakest precondition¹ of p with respect to the assignment part of the operation $\tau_i(l_i, l'_i)$. Add the atomic propositions occurring in $wp_{\tau_i(l_i, l'_i)}(p)$ as new predicates to $Pred$.

We exemplify our iterative abstraction refinement approach based on the

¹Computable via an SMT solver with built-in linear integer arithmetic theory. In our approach we use Z3 [13].

1 simple system Sys and the corresponding concrete control flow graph G_c de-
 2 picted in Figure 6.




Figure 6: Concurrent system Sys and corresponding concrete control flow graph G_c .

3 Here we have a single process operating on the integer variable y and we
 4 want to check whether there exists an execution that finally reaches control flow
 5 location 1. Thus, the temporal logic property of interest is $\mathbf{F}(loc = 1)$. In the
 6 first iteration, we start with bound $b = 1$ and $Pred = \emptyset$. The corresponding
 7 abstract control flow graph G_{a_1} , computable with the abstractor 3Spot [2], is
 8 depicted in Figure 7.




Figure 7: Abstract control flow graph G_{a_1} .

In order to solve the corresponding three-valued bounded model checking problem $[M(G_{a_1}) \models_{E,1} \mathbf{F}(loc = 1)]$ we construct the propositional logic encoding that now comprises the causes of uncertainty:

$$\begin{aligned} & \llbracket Sys(G_{a_1}), \mathbf{F}(loc = 1) \rrbracket_1 = \\ & \underbrace{(\neg l_0) \wedge ((\neg l_0 \wedge \neg l_1 \wedge \perp^{(0,0)_0}) \vee (\neg l_0 \wedge l_1 \wedge \perp^{(0,1)_0}) \vee (l_0 \wedge l_1))}_{Trans_{0,1}} \wedge \underbrace{(l_0 \vee l_1)}_{\llbracket \mathbf{F}(loc=1) \rrbracket_1^0} \end{aligned}$$

$\underbrace{(\neg l_0)}_{Init_0}$

Since our system consists of a single process and only one digit is necessary to encode the control flow, we can omit the process- and digit-indices of the atoms. Solely the position index is required for the encoding. Now we run the associated satisfiability tests. We get

$$\text{SAT}(\llbracket Sys(G_{a_1}), \mathbf{F}(loc = 1) \rrbracket_1 [\perp \mapsto true]) = true$$

and the corresponding satisfying truth assignment $\alpha : l_0 \mapsto false, l_1 \mapsto true$. Moreover, we get

$$\text{SAT}(\llbracket Sys(G_{a_1}), \mathbf{F}(loc = 1) \rrbracket_1 [\perp \mapsto false]) = false$$

Hence, $[M(G_{a_1}) \models_{E,1} \mathbf{F}(loc = 1)]$ yields *unknown* and α characterises an unconfirmed witness path $\langle 0 \rangle \xrightarrow{\perp} \langle 0 \rangle$ in the abstract state space where $\xrightarrow{\perp}$ denotes an *unknown* transition between the states. Next we apply the procedure *Cause-Guided Abstraction Refinement*. Remember that SAT solvers always operate on formulae transferred into conjunctive normal form (CNF). The current encoding is equivalent to the following formula in CNF²

$$\begin{aligned} & (\neg l_0) \wedge (l_0 \vee \neg l_1 \vee \perp_{(0,1)_0}) \wedge (l_1 \vee \perp_{(0,0)_0}) \\ & \wedge (l_0 \vee \perp_{(0,0)_0} \vee \perp_{(0,1)_0}) \wedge (l_1 \vee \perp_{(0,0)_0} \vee \perp_{(0,1)_0}) \wedge (l_0 \vee l_1) \end{aligned}$$

Under the assignment $\alpha[\perp \mapsto false]$ (the assignment α extended by the assignment that maps all \perp 's to *false*) we get the following set of unsatisfied clauses for our encoding:

$$U = \{(l_0 \vee \neg l_1 \vee \perp_{(0,1)_0}), (l_0 \vee \perp_{(0,0)_0} \vee \perp_{(0,1)_0})\}$$

A corresponding set of causes of uncertainty that covers U is

$$Causes = \{(0, 1)_0\}$$

1 since $\perp_{(0,1)_0}$ occurs in all clauses of U . $(0, 1)_0$ indicates that at the current
2 level of abstraction uncertainty is caused by the missing guard of the operation
3 $\tau(0, 1)$. We have that $\tau(0, 1) = \neg(y > 0)$ in the concrete system. Hence, we add
4 $(y > 0)$ to the set of predicates: $Pred := Pred \cup \{(y > 0)\}$ and proceed with the
5 next iteration.

6 In the second iteration, we have $b = 1$ and $Pred = \{(y > 0)\}$. The corre-
7 sponding abstract control flow graph G_{a_2} is depicted in Figure 8.

In order to solve $[M(G_{a_2}) \models_{E,1} \mathbf{F}(loc = 1)]$ we construct the following encoding ($(y > 0)$ abbreviated by p):

$$\begin{aligned} & \llbracket Sys(G_{a_2}), \mathbf{F}(loc = 1) \rrbracket_1 = \\ & (\neg l_0 \wedge \neg p[u]_0 \wedge p[t]_0) \wedge ((\neg l_0 \wedge \neg l_1 \wedge enc(\tau(0, 0))_{0,1}) \\ & \vee (\neg l_0 \wedge l_1 \wedge enc(\tau(0, 1))_{0,1}) \vee (l_0 \wedge l_1 \wedge enc(\tau(1, 1))_{0,1})) \wedge (l_0 \vee l_1) \end{aligned}$$

²For the sake of simplicity we use a standard CNF transformation in this illustrating example. Note that in our implementation we use the more compact Tseitin CNF transformation which introduces additional auxiliary atoms. Hence, we would get a slightly different CNF formula and unsatisfied clauses. Nevertheless, these clauses would hint at exactly the same causes of uncertainty.

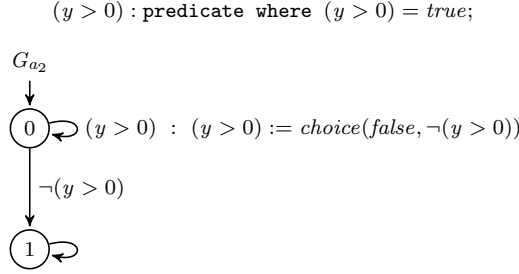


Figure 8: Abstract control flow graph G_{a_2} .

with

$$\begin{aligned}
& enc(\tau(0, 0))_{0,1} = \\
& ((p[u]_0 \wedge \perp^{\mathbf{po}}) \vee (\neg p[u]_0 \wedge p[t]_0)) \\
& \wedge (((p[u]_0 \wedge \perp^{\mathbf{po}}) \vee (\neg p[u]_0 \wedge \neg p[t]_0)) \wedge (\neg p[u]_1 \wedge \neg p[t]_1)) \\
& \vee ((p[u]_0 \vee (\neg p[u]_0 \wedge p[t]_0)) \wedge (p[u]_1))
\end{aligned}$$

and $\tau(0, 1)_{0,1}$ and $\tau(1, 1)_{0,1}$ encoded analogously. As we can see, uncertainty is now potentially caused by predicate p evaluating to \perp at position 0. Now we run the associated satisfiability tests. We get

$$\text{SAT}(\llbracket Sys(G_{a_2}), \mathbf{F}(loc = 1) \rrbracket_1[\perp \mapsto true]) = false$$

and

$$\text{SAT}(\llbracket Sys(G_{a_2}), \mathbf{F}(loc = 1) \rrbracket_1[\perp \mapsto false]) = false$$

- 1 Hence, $[M(G_{a_2}) \models_{E,1} \mathbf{F}(loc = 1)]$ yields *false*, which indicates that there
- 2 does not exist a 1-bounded witness path for $\mathbf{F}(loc = 1)$. Consequently, we
- 3 increment the bound: $b := b + 1$ and proceed with the next iteration.

In the third iteration, we have $b = 2$ and still $Pred = \{(y > 0)\}$. Hence, we continue with the abstract the control flow graph G_{a_2} . In order to solve $[M(G_{a_2}) \models_{E,2} \mathbf{F}(loc = 1)]$ we construct the following encoding:

$$\begin{aligned}
& \llbracket Sys(G_{a_2}), \mathbf{F}(loc = 1) \rrbracket_2 = \\
& (\neg l_0 \wedge \neg p[u]_0 \wedge p[t]_0) \wedge \bigwedge_{k=0}^1 ((\neg l_k \wedge \neg l_{k+1} \wedge enc(\tau(0, 0))_{k,k+1}) \\
& \vee (\neg l_k \wedge l_{k+1} \wedge enc(\tau(0, 1))_{k,k+1}) \vee (l_k \wedge l_{k+1} \wedge enc(\tau(1, 1))_{k,k+1})) \\
& \wedge (l_0 \vee l_1 \vee l_2)
\end{aligned}$$

Now we run the associated satisfiability tests. We get

$$\text{SAT}(\llbracket Sys(G_{a_2}), \mathbf{F}(loc = 1) \rrbracket_2[\perp \mapsto true]) = true$$

and the corresponding satisfying truth assignment $\alpha : l_0 \mapsto false, l_1 \mapsto false, l_2 \mapsto true, p[u]_0 \mapsto false, p[t]_0 \mapsto true, p[u]_1 \mapsto true, p[t]_1 \mapsto true, p[u]_2 \mapsto true, p[t]_2 \mapsto true$. Moreover, we get

$$\text{SAT}(\llbracket Sys(G_{a_2}), \mathbf{F}(loc = 1) \rrbracket_2[\perp \mapsto false]) = false$$

Hence, $[M(G_{a_2}) \models_{E,2} \mathbf{F}(loc = 1)]$ yields *unknown* and α characterises an unconfirmed witness path $\langle 0, p = true \rangle \rightarrow \langle 0, p = \perp \rangle \xrightarrow{\perp} \langle 1, p = \perp \rangle$ in the abstract state space. Next we apply the procedure *Cause-Guided Abstraction Refinement*. After deriving the set U of clauses of $\llbracket Sys(G_{a_2}), \mathbf{F}(loc = 1) \rrbracket_2$ that are unsatisfied under the assignment $\alpha[\perp \mapsto false]$, we determine a corresponding set of causes covering U . We get:

$$Causes = \{p_1\}$$

1 p_1 indicates that at the current level of abstraction uncertainty is caused
2 by the predicate p evaluating to *unknown* at position 1 of the witness path
3 characterised by α . Now we determine the last predecessor position where
4 the value of p is known, that is the greatest $k < 1$ with $\alpha(p[u]_k) = false$.
5 This holds for $k = 0$. Hence, the transition from position 0 to 1 along the
6 witness path characterised by α makes p *unknown*. We have that $\alpha(l_0) = false$
7 and $\alpha(l_1) = false$, which indicates that the transition from position 0 to 1 is
8 associated with the operation $\tau(0, 0)$. From the concrete control flow graph
9 we get that the assignment part of $\tau(0, 0)$ is $y := y - 1$. Thus, the weakest
10 precondition of $p = (y > 0)$ with respect to $\tau(0, 0)$ is $wp_{\tau(0,0)}(y > 0) = (y >$
11 $0)[y/y - 1] = (y - 1 > 0) = (y > 1)$. Hence, we add $(y > 1)$ to the set of
12 predicates and proceed with the next iteration.

13 In the forth iteration, we have $b = 2$ and $Pred = \{(y > 0), (y > 1)\}$. The
14 corresponding abstract control flow graph G_{a_3} is depicted in Figure 9.

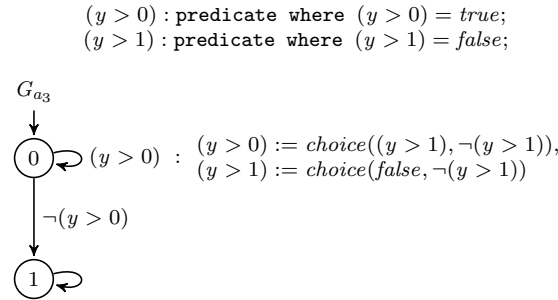


Figure 9: Abstract control flow graph G_{a_3} .

In order to solve $[M(G_{a_3}) \models_{E,2} \mathbf{F}(loc = 1)]$ we construct the encoding $\llbracket Sys(G_{a_3}), \mathbf{F}(loc = 1) \rrbracket_2$ and run the associated satisfiability tests. We get

$$\text{SAT}(\llbracket Sys(G_{a_3}), \mathbf{F}(loc = 1) \rrbracket_2[\perp \mapsto true]) = true$$

and

$$\text{SAT}(\llbracket \text{Sys}(G_{a_3}), \mathbf{F}(\text{loc} = 1) \rrbracket_2[\perp \mapsto \text{false}]) = \text{true}$$

1 and the corresponding satisfying truth assignment $\alpha : l_0 \mapsto \text{false}, l_1 \mapsto \text{false},$
 2 $l_2 \mapsto \text{true}, p[u]_0 \mapsto \text{false}, p[t]_0 \mapsto \text{true}, p[u]_1 \mapsto \text{false}, p[t]_1 \mapsto \text{false}, p[u]_2 \mapsto$
 3 $\text{false}, p[t]_2 \mapsto \text{false}, q[u]_0 \mapsto \text{false}, q[t]_0 \mapsto \text{false}, q[u]_1 \mapsto \text{true}, q[t]_1 \mapsto \text{false},$
 4 $q[u]_2 \mapsto \text{true}, q[t]_2 \mapsto \text{false}$ where p abbreviates $(y > 0)$ and q abbreviates
 5 $(y > 1)$. We can immediately conclude that $[M(G_{a_3}) \models_{E,2} \mathbf{F}(\text{loc} = 1)]$ yields
 6 *true*, which indicates that α characterises a definite 2-bounded witness path
 7 $\langle 0, p = \text{true}, q = \text{false} \rangle \rightarrow \langle 0, p = \text{false}, q = \perp \rangle \rightarrow \langle 1, p = \text{false}, q = \perp \rangle$
 8 for $\mathbf{F}(\text{loc} = 1)$. This outcome completes our verification task. Within four
 9 iterations of cause-guided abstraction refinement resp. bound incrementation
 10 we have automatically proven that the property of interest holds for the system.

11 Thus, given a software verification task to be solved, our cause-guided refine-
 12 ment approach enables us to automatically reach the right level of abstraction
 13 in order to obtain a definite result in verification. Next, we present the imple-
 14 mentation of our encoding-based model checking technique and we report on
 15 experimental results.

16 7. Implementation and Experiments

17 In this section we introduce the implementation of our approach and we
 18 outline extensions based on previous work on spotlight abstraction [2] and
 19 symmetry-based parameterised verification [14]. Moreover, we present experi-
 20 mental results. We have implemented a SAT-based bounded model checker for
 21 three-valued abstractions of concurrent software systems.³ Our tool employs
 22 the abstractor 3Spot [2] that builds abstract control flow graphs for a given
 23 concurrent system *Sys* and a set of predicates *Pred*. 3Spot supports almost all
 24 control structures of the C language as well as *int*, *bool* and *semaphore* as data
 25 types. Based on the CFGs and an input LTL formula ψ , our tool automatically
 26 constructs an encoding $\llbracket \text{Sys}, \psi \rrbracket_b$ of the corresponding verification task. The
 27 tool iteratively refines the abstraction in case of an *unknown* result and incre-
 28 ments the bound in case of a *false* result. It terminates once a *true* result is
 29 obtained or *false* result is obtained for a predefined threshold of the bound: In
 30 each iteration the two instances of the encoding are processed by a solver thread
 31 of the SAT solver Sat4j [15]. A *true* result for $\llbracket \text{Sys}, \psi \rrbracket_b[\perp \mapsto \text{false}]$ can be imme-
 32 diately transferred to the corresponding model checking problem $[M \models_{E,b} \psi]$.
 33 The same holds for a *false* result for $\llbracket \text{Sys}, \psi \rrbracket_b[\perp \mapsto \text{true}]$ if b represents a com-
 34 pleteness threshold of the verification task [11]. In case of an *unknown* result
 35 we apply cause-guided abstraction refinement as defined in the previous section.
 36 For *true* and *unknown* results, we additionally output a definite resp. uncon-
 37 firmed witness path for the property ψ in the form of an assignment satisfying

³available at www.github.com/ssfm-up/TVMC

$\llbracket Sys, \psi \rrbracket_b$. The tool chain of our model checker is depicted in Figure 10.

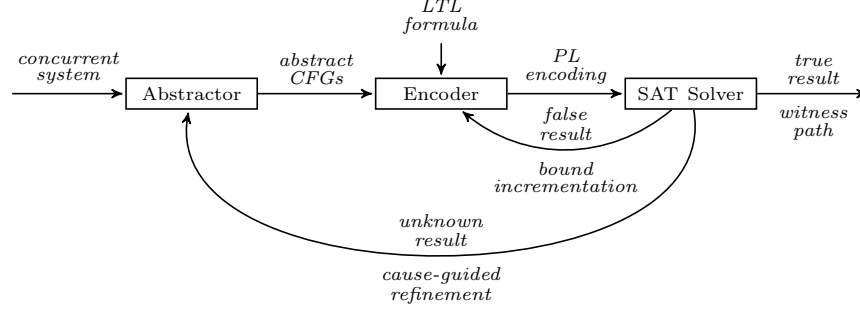


Figure 10: Tool chain.

We now illustrate how our tool systematically solves verification tasks via three-valued abstraction and cause-guided refinement. An example system $Sys = \parallel_{i=1}^n P_i$ implementing a solution to the dining philosophers problem is depicted in Figure 11. Here we have $n \in \mathbb{N}$ philosopher processes and the same number of binary semaphore variables modelling the forks. Processes P_i with $i < n$ continuously attempt to first acquire the semaphore y_i and second y_{i+1} , whereas process P_n attempts to acquire first y_n and then y_1 . Once a process has successfully acquired both semaphores it consecutively releases them and attempts to acquire them again.




Figure 11: Dining philosophers system Sys .

Our tool generally searches for violations of desirable properties. For an instantiation of the dining philosophers system with $n = 2$ the violation of mutual exclusion can be expressed in LTL as

$$\psi = \mathbf{F}((loc_1 = 10) \wedge (loc_2 = 10)).$$

Hence, we want to check whether the processes P_1 and P_2 will be ever at their critical location 10 at the same time. Starting with $Pred = \emptyset$ and $b = 1$ our tool automatically constructs the corresponding abstract control flow graphs and the encoding $\llbracket Sys, \psi \rrbracket_b$. Next, it iteratively increments the bound and refines the initial abstraction. For our example the bound will be increased until an unconfirmed witness path for the property of interest will be detected

at $b = 4$:

$$\langle 00, 00 \rangle \xrightarrow{\perp} \langle 01, 00 \rangle \xrightarrow{\perp} \langle 01, 01 \rangle \xrightarrow{\perp} \langle 10, 01 \rangle \xrightarrow{\perp} \langle 10, 10 \rangle$$

1 Then cause-guided refinement will add the predicates y_1 and y_2 in a single
2 step, which yields the level of abstraction characterised by the abstract control
3 flow graphs depicted in Figure 12. Finally the bound will be further increased
4 until a completeness threshold is reached, which is the case for $b = 64$ for this
5 verification task. – A technique for computing over-approximations of complete-
6 ness thresholds is introduced in [11]. Completeness thresholds for checking LTL
7 properties that are restricted to the temporal operators **F** and **G** are linear in
8 the size of the abstraction, i.e. in the number of abstract states.

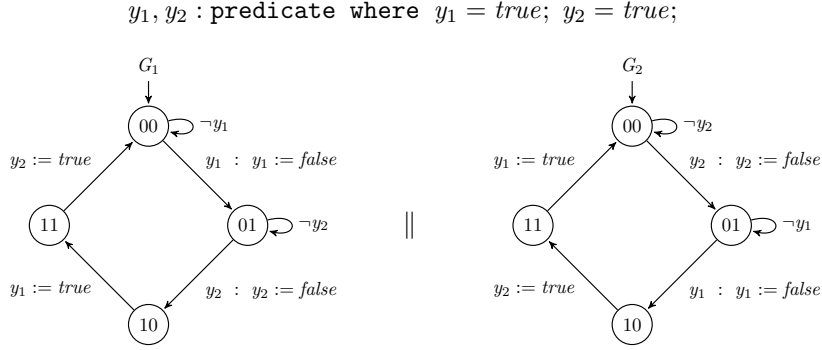


Figure 12: Abstraction of the dining philosophers system with $n = 2$.

9 Via SAT solving we obtain a *false* result in the final iteration, which allows
10 us to conclude that mutual exclusion is not violated for the instantiation of the
11 dining philosophers system with 2 processes.

The abstractor 3Spot does not only allow us to abstract fixed-sized systems, it also allows to construct finite abstractions of parameterised systems with an unbounded number of uniform processes [14]. For this 3Spot combines three-valued predicate abstraction with *spotlight abstraction* [2]. The basic concept of spotlight abstraction is to partition a parallel composition of processes into a *spotlight* and a *shade*. The control flow of spotlight processes is then explicitly considered under abstraction, whereas the processes in the shade get summarised into a single abstract process P_{\perp} that approximates their behaviour with regard to three-valued logic. Hence, predicates over variables that are modified by processes in the shade may be set to *unknown* by P_{\perp} . For our example verification task, spotlight abstraction can be applied to the process partition $\text{Spotlight} = \{P_1, P_2\}$ and $\text{Shade} = \{P_3, \dots, P_n\}$, i.e. we consider P_1 and P_2 explicitly and summarise the parameterised number of processes P_3 to P_n into P_{\perp} . The property ψ can also be disproven for the abstraction $P_1 \parallel P_2 \parallel P_{\perp}$ with our tool, which allows us to conclude that mutual exclusion (with regard to P_1 and P_2) is not violated for all instantiations of the dining philosopher system with $n \geq 2$ processes. The LTL formula ψ characterises a local property

since it refers to particular processes of a parameterised system. However, as shown in [14], symmetry arguments enable us to transfer this result to arbitrary pairs of processes in the system. We can conclude that

$$\psi_{global} = \exists 1 \leq i, j \leq n, i \neq j : \mathbf{F}((loc_i = 10) \wedge (loc_j = 10))$$

1 does not hold for any instantiation of the dining philosophers system, i.e. no pair
 2 of processes will be ever at their critical location at the same time. More details
 3 about spotlight abstraction and its combination with symmetry arguments can
 4 be found in [2] and [14].

A distinct feature of our approach is the verification of liveness properties of concurrent systems under fairness assumptions. The formula

$$\psi' = \bigvee_{i=1}^n \mathbf{F}(\mathbf{G}\neg(loc_i = 10))$$

characterises the violation of a liveness property regarding our example system. It states that eventually some philosopher process will nevermore reach its critical location. For the fix-sized instantiation with $n = 2$ processes and starting with $Pred = \{progress_1, progress_2\}$ (fairness predicates only) and $b = 1$ our tool constructs the encoding $\llbracket Sys, fair \wedge \psi' \rrbracket_b$. Within three iterations over b and a refinement step adding the predicates y_1 and y_2 we can already detect a satisfying assignment for the encoding that characterises a fair path with a loop where both P_1 and P_2 reach location 01 and remain there forever:

$$\begin{array}{c} \langle 00, 00, y_1 = true, y_2 = true, progress_1 = false, progress_2 = false \rangle \\ \downarrow \\ \langle 01, 00, y_1 = false, y_2 = true, progress_1 = true, progress_2 = false \rangle \\ \downarrow \\ \langle 01, 01, y_1 = false, y_2 = false, progress_1 = false, progress_2 = true \rangle \\ \downarrow \uparrow \\ \langle 01, 01, y_1 = false, y_2 = false, progress_1 = true, progress_2 = false \rangle \end{array}$$

5 Thus, we have proven that liveness of Sys is violated under fairness. With
 6 our tool we could also successfully detect liveness violations in generalisations
 7 of the dining philosophers system with more philosophers and semaphores. The
 8 experimental results with regard to the formula ψ' are depicted in Table 1.
 9 The experiments were conducted on a 1.6 GHz Intel Core i7 system with 8 GB
 10 memory. We measured the final bound, the number of refinement steps, the
 11 final number of predicates as well as the overall time for encoding and SAT-
 12 based model checking in all iterations. Beside promising performance results we
 13 discovered that for all verification tasks our cause-guided refinement detected
 14 and added all necessary predicates within a single step. Thus, uncertainty in
 15 the abstraction due missing predicates was ruled out promptly.

16 The faultiness in terms of liveness of the example system can easily corrected
 17 by by changing the order of requests of the the n -th philosopher process from

philosophers	final bound	refinement steps	final number of predicates	overall time
2	3	1	4	1.13s
3	5	1	6	2.12s
4	7	1	8	4.69s
5	9	1	10	12.4s
6	11	1	12	38.1s
7	13	1	14	379s
8	15	1	16	75.0m

Table 1: Experimental results.

1 first y_n then y_1 to first y_1 then y_n . (Remark: Since spotlight abstraction in-
 2 herently abstracts away the *order* in which operations occurring in processes in
 3 the shade may be executed, this abstraction technique could not be utilised for
 4 these liveness verification tasks.) While bounded model checking is generally
 5 considered as a technique for error detection, we were also able to *prove* live-
 6 ness of instantiations of the corrected system, which required to let our tool run
 7 until the (linear) completeness threshold for the verification task was reached.
 8 The verification of the corrected system with 2 philosophers took 39.4 minutes
 9 until the bound reached the completeness threshold of 64. Although our ex-
 10 periments already showed encouraging results, we expect that we can further
 11 enhance the performance of our tool based on optimisations that we mention in
 12 the conclusion of this paper.

13 8. Related Work

14 Our SAT-based software verification technique is related to a number of
 15 existing approaches in the field of bounded model checking for software. The
 16 bounded model checker CBMC [16] supports the verification of *sequential C*
 17 *programs*. It is based on a Boolean abstraction of the input program and it
 18 allows for checking buffer overflows, pointer safety and assertions, but not full
 19 LTL properties. A similar tool is F-Soft [17]. This bounded model checker
 20 for sequential programs is restricted to the verification of *reachability proper-*
 21 *ties*. While CBMC and F-Soft support a wider range of program constructs like
 22 pointers and recursion, our approach focusses on the challenges associated with
 23 *concurrency* and the verification of *liveness properties* under fairness. The tool
 24 TCBMC [18] is an extension of CBMC for verifying *safety properties* of *concur-*
 25 *rent programs*. TCBMC introduces the concept of bounding context switches
 26 between processes, which is a special abstraction technique for reducing concu-
 27 rrency. Our approach supports the process summarisation abstraction of 3Spot
 28 [2], which allows us to reduce the complexity induced by concurrency in a dif-
 29 ferent way. The verification of concurrent C programs is also addressed in [19].
 30 The authors introduce a tool that translates C programs into a TLA+ [20]
 31 specification which is then model checked via an explicit-state approach.

32 In contrast to the above mentioned tools, we employ *three-valued abstraction*,
 33 which preserves *true* and *false* results in verification. *Three-valued bounded*

1 *model checking* is addressed in [7] and [8]. However, only in the context of
2 *hardware verification* [7] resp. assuming that an *explicit three-valued Kripke*
3 *structure* is given [8]. To the best of our knowledge, our approach is the first
4 that supports software verification under fairness via an immediate propositional
5 logic encoding and SAT-based BMC.

6 Abstraction refinement for SAT-based model checking is addressed in [21,
7 22, 23]. The hardware verification approach presented in [21] employs Boolean
8 abstraction by means of variable hiding. Counterexamples detected in the ab-
9 stract model are simulated on the concrete model via SAT solving. Unsatisfia-
10 bility results correspond to spurious counterexamples. In this case abstraction
11 refinement is applied by deriving new variables from the unsatisfiable core. Very
12 similar approaches are used in [22] and [23]. The authors of [22] additionally
13 show that their technique yields an over-approximative abstraction that pre-
14 serves safety and always has a completeness threshold for not only refuting but
15 also proving properties. The authors of [23] generate refinement interpolants
16 from simulated spurious counterexamples in SAT-based model checking. In our
17 three-valued approach we do not have to simulate abstract counterexamples.
18 In case of an *unknown* result, our path characterising assignment allows us to
19 derive new predicates from the set of unsatisfied clauses that is typically sig-
20 nificantly smaller than the unsatisfiable core. Unconfirmed witness paths in a
21 three-valued abstract model are also used for refinement in [2]. However, the
22 proposed approach requires to explicitly generate and analyse paths, whereas
23 our refinement happens based on unsatisfied clauses that implicitly represent
24 uncertainty in the abstraction.

25 9. Conclusion and Outlook

26 We introduced a verification technique for concurrent software systems based
27 on three-valued abstraction, cause-guided refinement and SAT-based bounded
28 model checking. We defined a direct propositional logic encoding of software
29 verification tasks and we proved that our encoding is sound in the sense that SAT
30 results can be straightforwardly transferred to the corresponding model checking
31 problem. Hence, the expensive construction and exploration of an explicit state
32 space model is not necessary. Our tool enables the verification of safety and
33 liveness properties under fairness. With cause-guided refinement we introduced
34 an automatic technique for systematically reaching the right level of abstraction
35 in order to obtain a definite outcome in verification. Refinement steps are
36 straightforwardly derived from clauses of the encoding that are unsatisfied under
37 assignments characterising potential witness/error paths. Due to the efficiency
38 of modern SAT solvers we achieve promising performance results with our overall
39 approach.

40 As future work we plan to experimentally evaluate our approach based on
41 case studies on concurrent software systems and compare our model checker
42 to similar tools. Moreover, we intend to optimise our technique by integrating
43 *incremental* SAT solving [24] into the tool and by developing a concept for
44 reusing parts of the encoding between the consecutive refinement iterations.

1 Finally, we want to develop SAT solving *heuristics* tailored to the structure of
 2 our encodings [25] in order to further accelerate our approach.

3 **Acknowledgements**

4 We thank Matthias Harvey for help with implementing our approach.

5 **References**

- 6 [1] S. Shoham, O. Grumberg, 3-valued abstraction: More precision at less cost,
 7 Information and Computation 206 (11) (2008) 1313–1333.
- 8 [2] J. Schrieb, H. Wehrheim, D. Wonisch, Three-valued spotlight abstractions,
 9 in: A. Cavalcanti, D. Dams (Eds.), FM, Vol. 5850 of LNCS, Springer, 2009,
 10 pp. 106–122.
- 11 [3] G. Bruns, P. Godefroid, Model checking partial state spaces with 3-valued
 12 temporal logics, in: CAV 1999, LNCS, Springer Berlin Heidelberg, 1999,
 13 pp. 274–287.
- 14 [4] N. Timm, H. Wehrheim, M. Czech, Heuristic-guided abstraction re-
 15 finement for concurrent systems, in: T. Aoki, K. Taguchi (Eds.),
 16 ICFEM, Vol. 7635 of LNCS, Springer, 2012, pp. 348–363. doi:10.1007/
 17 978-3-642-34281-3_25.
- 18 [5] A. Cimatti, E. Clarke, F. Giunchiglia, M. Roveri, NuSMV: a new symbolic
 19 model checker, Int. Jour. on Softw. Tools for Techn. Transfer 2 (4) (2000)
 20 410–425.
- 21 [6] A. Biere, A. Cimatti, E. M. Clarke, O. Strichman, Y. Zhu, Bounded model
 22 checking., Handbook of Satisfiability 185 (2009) 457–481.
- 23 [7] O. Grumberg, 3-Valued Abstraction for (Bounded) Model Checking,
 24 Springer Berlin Heidelberg, Berlin, Heidelberg, 2009, pp. 21–21. doi:
 25 10.1007/978-3-642-04761-9_2.
- 26 [8] H. Wehrheim, Bounded model checking for partial Kripke structures, in:
 27 J. Fitzgerald, A. Haxthausen (Eds.), ICTAC, Vol. 5160 of LNCS, Springer,
 28 2008, pp. 380–394. doi:10.1007/978-3-540-85762-4_26.
- 29 [9] N. Timm, S. Gruner, M. Harvey, A Bounded Model Checker for
 30 Three-Valued Abstractions of Concurrent Software Systems, Springer
 31 International Publishing, Cham, 2016, pp. 199–216. doi:10.1007/
 32 978-3-319-49815-7_12.
 33 URL http://dx.doi.org/10.1007/978-3-319-49815-7_12
- 34 [10] M. Fitting, Kleene’s 3-valued logics and their children, Fund. Inf. 20 (1-3)
 35 (1994) 113–131.

- 1 [11] D. Kroening, J. Ouaknine, O. Strichman, T. Wahl, J. Worrell, Linear com-
2 pleteness thresholds for bounded model checking, in: CAV, Springer, 2011,
3 pp. 557–572.
- 4 [12] N. Timm, Bounded model checking für partielle Systeme, Masters thesis,
5 University of Paderborn.
- 6 [13] L. Moura, N. Björner, Z3: An efficient SMT solver, in: C. R. Ramakrishnan,
7 J. Rehof (Eds.), Tools and Algorithms for the Construction and Analysis of
8 Systems, Vol. 4963 of Lecture Notes in Computer Science, Springer-Verlag
9 Berlin Heidelberg, 2008, pp. 337–340. doi:10.1007/978-3-540-78800-3_
10 24.
11 URL http://dx.doi.org/10.1007/978-3-540-78800-3_24
- 12 [14] N. Timm, H. Wehrheim, On symmetries and spotlights - verifying param-
13 eterised systems, in: J. Dong, H. Zhu (Eds.), ICFEM, Vol. 6447 of LNCS,
14 Springer, 2010, pp. 534–548.
- 15 [15] D. Le Berre, A. Parrain, The Sat4j library, release 2.2, Journal on Satisfi-
16 ability, Boolean Modeling and Computation 7 (2010) 59–64.
- 17 [16] D. Kroening, M. Tautschnig, CBMC-C bounded model checker, in:
18 TACAS, Springer, 2014, pp. 389–391.
- 19 [17] F. Ivančić, Z. Yang, M. K. Ganai, A. Gupta, I. Shlyakhter, P. Ashar, F-Soft:
20 Software verification platform, in: International Conference on Computer
21 Aided Verification, Springer, 2005, pp. 301–306.
- 22 [18] I. Rabinovitz, O. Grumberg, Bounded model checking of concurrent pro-
23 grams, in: CAV, Springer, 2005, pp. 82–97.
- 24 [19] A. Methni, M. Lemerre, B. Ben Hedia, S. Haddad, K. Barkaoui, Specifying
25 and Verifying Concurrent C Programs with TLA+, Springer International
26 Publishing, Cham, 2015, pp. 206–222. doi:10.1007/978-3-319-17581-2_
27 14.
28 URL http://dx.doi.org/10.1007/978-3-319-17581-2_14
- 29 [20] S. Merz, The Specification Language TLA+, Springer Berlin Hei-
30 delberg, Berlin, Heidelberg, 2008, pp. 401–451. doi:10.1007/
31 978-3-540-74107-7_8.
32 URL http://dx.doi.org/10.1007/978-3-540-74107-7_8
- 33 [21] A. Gupta, O. Strichman, Abstraction Refinement for Bounded Model
34 Checking, Springer Berlin Heidelberg, Berlin, Heidelberg, 2005, pp. 112–
35 124. doi:10.1007/11513988_11.
36 URL http://dx.doi.org/10.1007/11513988_11
- 37 [22] D. Kroening, Computing over-approximations with bounded model check-
38 ing, Electron. Notes Theor. Comput. Sci. 144 (1) (2006) 79–92. doi:
39 10.1016/j.entcs.2005.07.021.
40 URL <http://dx.doi.org/10.1016/j.entcs.2005.07.021>

- 1 [23] C. Y. Wu, C. A. Wu, C.-Y. Lai, C. Y. Huang, A counterexample-guided
2 interpolant generation algorithm for sat-based model checking, in: 2013
3 50th ACM/EDAC/IEEE Design Automation Conference (DAC), 2013, pp.
4 1–6.
- 5 [24] A. Nadel, V. Ryvchin, O. Strichman, Ultimately incremental sat, in: Inter-
6 national Conference on Theory and Applications of Satisfiability Testing,
7 Springer, 2014, pp. 206–218.
- 8 [25] N. Timm, S. Gruner, P. Sibanda, Model checking of concurrent software
9 systems via heuristic-guided sat solving, in: International Conference on
10 Fundamentals of Software Engineering, Springer, 2017.